

Recommender Systems

User-Based Collaborative Filtering


Rodrygo L. T. Santos
rodrygo@dcc.ufmg.br

The recommendation problem




$$f(u, i)$$

The recommendation problem



The recommendation problem




The recommendation problem




$$f(u, i)$$

Cold-start user




Non-personalized recommendation



most recent
most popular
best rated


**What if we
know the
user?**

Personalized recommendation




exploit **paths** via
other users for
personalization

Personalized recommendation




**more users,
better signal?**

Personalized recommendation




**coincident users,
better signal?**

Personalized recommendation



**similar users,
better signal!**

User-based collaborative filtering



User-based collaborative filtering

	i_1	i_2	i_3	i_4
u_1	+1	+1	+1	
u_2	+1			+1
u_3				+1


unary feedback
(e.g., click)

User-based collaborative filtering

	i_1	i_2	i_3	i_4
u_1	+1	+1	+1	
u_2	+1			+1
u_3		-1	-1	+1

binary feedback
(e.g., like / dislike)

User-based collaborative filtering

	i_1	i_2	i_3	i_4
u_1	5	3	3	
u_2	5			3
u_3		1	2	1

graded feedback
(e.g., rate)

Breaking it down

Normalizing ratings

Computing similarities

Selecting neighborhoods

Aggregating ratings

Normalizing ratings

Users rate differently

- Some rate high, others low
- Some use more of the scale than others

Aggregation ignores these differences

- Normalization compensates for them


Mean-centering normalization

	i_1	i_2	i_3	i_4
u_1	5	3	3	
u_2	5			3
u_3		1	2	1

$$\bar{r}_{u_1} = 3.7$$

$$\bar{r}_{u_2} = 4.0$$

$$\bar{r}_{u_3} = 1.3$$



Mean-centering normalization

	i_1	i_2	i_3	i_4
u_1	1.3	-0.7	-0.7	
u_2	1.0			-1.0
u_3		-0.3	0.7	-0.3

Selecting neighborhoods

A few options

- All the neighbors
- Random neighbors
- All neighbors above a **similarity** threshold
- Top- k neighbors ranked by **similarity**

Computing similarities

In principle, any similarity would do

- Jaccard, Dice index
- Pearson correlation
- Cosine

Empirically, some perform better

Pearson correlation

$$s_{\vec{u}\vec{v}} = \frac{\text{cov}(\vec{u}\vec{v})}{\sigma_{\vec{u}}\sigma_{\vec{v}}} \\ \approx \frac{\sum_{i \in I_{uv}} (r_{ui} - \bar{r}_u)(r_{vi} - \bar{r}_v)}{\sqrt{\sum_{i \in I_{uv}} (r_{ui} - \bar{r}_u)^2} \sqrt{\sum_{i \in I_{uv}} (r_{vi} - \bar{r}_v)^2}}$$

- Built-in mean-centering normalization
- Computed over ratings in common (I_{uv})

What about little data?


Two users with one common rating

- $|I_{uv}| = 1 \rightarrow s_{\vec{u}\vec{v}} = 1$

Solution: weight similarity by confidence


- Simple approach: multiply by $\min(|I_{uv}|, 50)/50$
 - $|I_{uv}| < 50$: similarity scaled down by $|I_{uv}|/50$
 - $|I_{uv}| \geq 50$: similarity kept unscaled

Cosine similarity



From angles to cosines

Cosine is a monotonically decreasing function of the angle for the interval $[0^\circ, 180^\circ]$




*smaller the angle,
larger the cosine*

vs.


*larger the angle,
smaller the cosine*

Cosine similarity



$$\begin{aligned} s_{\vec{u}\vec{v}} &= \cos(\theta_{\vec{u}\vec{v}}) \\ &= \frac{\vec{u} \cdot \vec{v}}{\|\vec{u}\| \|\vec{v}\|} \\ &= \frac{\sum_{i \in I} r_{ui} r_{vi}}{\sqrt{\sum_{i \in I} r_{ui}^2} \sqrt{\sum_{i \in I} r_{vi}^2}} \end{aligned}$$

Cosine similarity



$$s_{\vec{u}\vec{v}} = \frac{4 \times 5 + 4 \times 2}{\sqrt{4^2 + 4^2} \sqrt{5^2 + 2^2}} \\ \approx 0.92$$

Cosine similarity

	i_1	i_2	i_3	i_4
u_1	1.3	-0.7	-0.7	
u_2	1.0			-1.0
u_3		-0.3	0.7	-0.3

$$s_{\vec{u}_1 \vec{u}_2} = +0.58$$

$$s_{\vec{u}_1 \vec{u}_3} = -0.17$$

Cosine similarity

$$s_{\vec{u}\vec{v}} = \cos(\theta_{\vec{u}\vec{v}}) = \frac{\vec{u} \cdot \vec{v}}{\|\vec{u}\| \|\vec{v}\|} = \frac{\sum_{i \in I} r_{ui} r_{vi}}{\sqrt{\sum_{i \in I} r_{ui}^2} \sqrt{\sum_{i \in I} r_{vi}^2}}$$

- In general: $s_{\vec{u}\vec{v}} \in [-1, 1]$
- With non-negative ratings: $s_{\vec{u}\vec{v}} \in [0, 1]$
- With mean-centering (aka ***adjusted cosine***)
 - Equivalent to Pearson... ***almost!***

Cosine vs. Pearson

Cosine has built-in significance weighting

- Similarity scaled by ratio

$$|I_u \cap I_v| / |I_u||I_v|$$

- Similar effect can be obtained by using **overall σ** instead of just **σ over common ratings** in Pearson

How many neighbors?

In theory, the more the better...

- ... if you have a good similarity metric
- Computational cost is also higher

In practice

- More neighbors → more noise
- Fewer neighbors → lower coverage

Aggregating ratings

A few options

- Min / max / average / median rating
- Weighted average (by similarity)
- Supervised aggregation

Common practice

- Weighted average: simple and effective

Aggregating ratings

	i_1	i_2	i_3	i_4	
u_1	1.3	-0.7	-0.7	?	$\hat{r}_{u_1 i_4} - \bar{r}_{u_1} = \tilde{r}_{u_1 i_4}$
u_2	1.0			-1.0	$r_{u_2 i_4} - \bar{r}_{u_2} = \tilde{r}_{u_2 i_4}$
u_3		-0.3	0.7	-0.3	$r_{u_3 i_4} - \bar{r}_{u_3} = \tilde{r}_{u_3 i_4}$

Aggregating ratings

	i_1	i_2	i_3	i_4
u_1	1.3	-0.7	-0.7	?
u_2	1.0			-1.0
u_3		-0.3	0.7	-0.3

$\tilde{r}_{u_1 i_4}$

$= \tilde{r}_{u_2 i_4}$

$= \tilde{r}_{u_3 i_4}$

$s_{\vec{u}_1 \vec{u}_2} = +0.58$

$s_{\vec{u}_1 \vec{u}_3} = -0.17$

Aggregating ratings

	i_1	i_2	i_3	i_4
u_1	1.3	-0.7	-0.7	?
u_2	1.0			-1.0
u_3		-0.3	0.7	-0.3

$$\begin{aligned}
 \tilde{r}_{u_1 i_4} &= \frac{\sum_{c=1}^k s_{\vec{u}_1 \vec{u}_c} \tilde{r}_{u_c i_4}}{\sum_{c=1}^k |s_{\vec{u}_1 \vec{u}_c}|} \\
 &= \frac{(0.58 \times -1.0) + (-0.17 \times -0.3)}{|-0.58| + |-0.17|} \\
 &= -0.71
 \end{aligned}$$

Aggregating ratings

	i_1	i_2	i_3	i_4
u_1	1.3	-0.7	-0.7	?
u_2	1.0			-1.0
u_3		-0.3	0.7	-0.3

$$\tilde{r}_{u_1 i_4} = -0.71$$

$$\begin{aligned}\hat{r}_{u_1 i_4} &= \tilde{r}_{u_1 i_4} + \bar{r}_{u_1} \\ &= -0.71 + 3.7 \\ &= 2.99\end{aligned}$$

Suggested configuration


Reasonable starting point

- Top- k neighbors ($k \approx 30$)
- Weighted averaging of scores
- Mean-centering normalization
- Cosine similarity

Optimal configuration is application-dependent

**How
efficient
is this?**

User-based collaborative filtering

	i_1	i_2	i_3	i_4
u_1	1.3	-0.7	-0.7	
u_2	1.0			-1.0
u_3		-0.3	0.7	-0.3

m users \times n items

Cosine between two users: $O(n)$

Number of potential neighbors: $O(m)$

Cost per user: $O(mn)$

User-based collaborative filtering

	u_1	u_2	u_3
u_1	$S_{\vec{u}_1 \vec{u}_1}$	$S_{\vec{u}_1 \vec{u}_2}$	$S_{\vec{u}_1 \vec{u}_3}$
u_2	$S_{\vec{u}_2 \vec{u}_1}$	$S_{\vec{u}_2 \vec{u}_2}$	$S_{\vec{u}_2 \vec{u}_3}$
u_3	$S_{\vec{u}_3 \vec{u}_1}$	$S_{\vec{u}_3 \vec{u}_2}$	$S_{\vec{u}_3 \vec{u}_3}$

m users \times n items

Cosine between two users: $O(n)$

Number of potential neighbors: $O(m)$

Cost per user: $O(mn)$

Cost for all users: $O(m^2n)$

Problem: m and n in the order of 10^7

Some simple optimizations

Can exploit matrix symmetry

- Only need to compute one direction

Can exploit matrix sparsity

- Only need to consider users with a common item

Still a costly operation to perform online

- ***Should we precompute similarities?***

Summary

User-based CF is simple and effective

- The oldest CF approach

Lots of configuration knobs

- Similarity functions, neighborhood selection, aggregation functions, rating normalization

Neighborhood selection is a bottleneck